

X


<https://swayam.gov.in>

https://swayam.gov.in/nc_details/NPTEL

harshaldharpure9922@gmail.com ▾

NPTEL (<https://swayam.gov.in/explorer?ncCode=NPTEL>) » Introduction to Large Language Models (LLMs)
(course)



Click to register
for Certification
exam

(https://examform.nptel.ac.in/2025_01/exam_form/dashboard)

If already
registered, click
to check your
payment status

Course outline

About NPTEL
()

How does an
NPTEL online
course work?
()

Week 1 ()

Week 2 ()

Week 3 ()

Week 8 : Assignment 8

The due date for submitting this assignment has passed.

Due on 2025-03-19, 23:59 IST.

Assignment submitted on 2025-03-13, 14:46 IST

1) Which factors influence the effectiveness of instruction tuning?

1 point

- ☒ The number of instruction templates used in training.
- ☐ The tokenization algorithm used by the model.
- ☒ The diversity of tasks in the fine-tuning dataset.
- ☒ The order in which tasks are presented during fine-tuning.

Yes, the answer is correct.

Score: 1

Accepted Answers:

The number of instruction templates used in training.

The diversity of tasks in the fine-tuning dataset.

The order in which tasks are presented during fine-tuning.

2) What are key challenges of soft prompts in prompt-based learning?

1 point

- ☐ Forward pass with them is computationally inefficient compared to that with hard prompts.
- ☒ They require additional training, unlike discrete prompts.
- ☒ They cannot be interpreted or used effectively by non-expert users.
- ☐ They require specialized architectures that differ from standard transformers.

Week 4 ()**Week 5 ()****Week 6 ()****Week 7 ()****Week 8 ()**

☐ Lec 21 :
Instruction
Tuning (unit?
unit=70&lesson
=71)

☐ Lec 22 :
Prompt-based
Learning (unit?
unit=70&lesson
=72)

☐ Lec 23 :
Advanced
Prompting and
Prompt
Sensitivity
(unit?
unit=70&lesson
=73)

☐ Lec 24 :
Alignment of
Language
Models-I (unit?
unit=70&lesson
=74)

☐ Lec 25 :
Alignment of
Language
Models-II (unit?
unit=70&lesson
=75)

☒ Lecture
Material (unit?
unit=70&lesson
=82)

☐ Feedback Form
(unit?)

Yes, the answer is correct.

Score: 1

Accepted Answers:

They require additional training, unlike discrete prompts.

They cannot be interpreted or used effectively by non-expert users.

3) Which statement best describes the impact of fine-tuning versus prompting in LLMs? **1 point**

- ☐ Fine-tuning is always superior to prompting in generalization tasks.
- ☐ Prompting requires gradient updates, while fine-tuning does not.
- ☒ Fine-tuning modifies the model weights permanently, while prompting does not.
- ☐ Prompting performs better on in-domain tasks compared to fine-tuning.

Yes, the answer is correct.

Score: 1

Accepted Answers:

Fine-tuning modifies the model weights permanently, while prompting does not.

4) Which of the following aspects of the model outputs are captured by POSIX? **1 point**

- ☒ Diversity in the responses to intent-preserving prompt variations
- ☐ Entropy of the distribution of response frequencies
- ☐ Time required to generate responses for intent-preserving prompt variations
- ☒ Variance in the log-likelihood of the same response for different input prompt variations

Partially Correct.

Score: 0.67

Accepted Answers:

Diversity in the responses to intent-preserving prompt variations

Entropy of the distribution of response frequencies

Variance in the log-likelihood of the same response for different input prompt variations

5) Which key mechanism makes Tree-of-Thought (ToT) prompting more effective than Chain-of-Thought (CoT)? **1 point**

- ☐ ToT uses reinforcement learning for better generalization.
- ☒ ToT allows backtracking to explore multiple reasoning paths.
- ☐ ToT reduces hallucination by using domain-specific heuristics.
- ☐ ToT eliminates the need for manual prompt engineering.

Yes, the answer is correct.

Score: 1

Accepted Answers:

ToT allows backtracking to explore multiple reasoning paths.

6) What is a key limitation of measuring accuracy alone when evaluating LLMs? **1 point**

- ☐ Accuracy is always correlated with model size.
- ☐ Accuracy cannot be measured on open-ended tasks.
- ☐ Accuracy is independent of the training dataset size.

unit=70&lesson=76)

**Quiz: Week 8 :
Assignment 8
(assessment?
name=77)**

Week 9 ()

Week 10 ()

Week 11 ()

Week 12 ()

**Year 2025
Solutions ()**

☒ Accuracy does not account for prompt sensitivity.

Yes, the answer is correct.

Score: 1

Accepted Answers:

Accuracy does not account for prompt sensitivity.

7) Why is instruction tuning not sufficient for aligning large language models?

1 point

- ☐ It does not generalize to unseen tasks.
- ☒ It cannot prevent models from generating undesired responses.
- ☐ It reduces model performance on downstream tasks.
- ☐ It makes models less capable of learning from new data.

Yes, the answer is correct.

Score: 1

Accepted Answers:

It cannot prevent models from generating undesired responses.

8) Why is KL divergence minimized in regularized reward maximization?

1 point

- ☐ To maximize the probability of generating high-reward responses.
- ☐ To make training more computationally efficient.
- ☐ To prevent the amplification of bias in training data.
- ☒ To ensure models do not diverge too far from the reference model.

Yes, the answer is correct.

Score: 1

Accepted Answers:

To ensure models do not diverge too far from the reference model.

9) What is the primary advantage of using the log-derivative trick in REINFORCE?

1 point

- ☐ Reducing data requirements
- ☐ Expanding the token vocabulary
- ☒ Simplifying gradient computation
- ☐ Improving sampling diversity

Yes, the answer is correct.

Score: 1

Accepted Answers:

Simplifying gradient computation

10) Which method combines reward maximization and minimizing KL divergence?

1 point

- ☐ REINFORCE
- ☐ Monte Carlo Approximation
- ☒ Proximal Policy Optimization
- ☐ Constitutional AI

Yes, the answer is correct.

Score: 1

Accepted Answers:

Proximal Policy Optimization