

# Quantization, Pruning and Distillation

Dinesh Raghu

Senior Researcher, IBM Research



**Introduction to Large Language Models**



# LLM Sizes




Image credits: <https://huggingface.co/blog/large-language-models>

# LLM Sizes






Image credits: <https://huggingface.co/blog/large-language-models>

# LLM Inference




Larger the model, larger the

1. GPU memory requirement
2. latency
3. inference cost
4. environmental concerns



Image credits: <https://huggingface.co/blog/large-language-models>

# LLM Inference

Image credits: openai.com





# LLM Inference

Calculate by:		Input tokens:	Output tokens:	Number of API calls:	
Tokens	Words	Characters			
		100	500	10000	
Provider	Model	Input price for 1M tokens	Output price for 1M tokens	Price per API call	Total price
 OpenAI	gpt-4o	\$5.00	\$15.00	\$0.0080	\$80.00
 OpenAI	gpt-4o-mini	\$0.15	\$0.60	\$0.0003	\$3.15

Why is gpt-4o-mini so cheap when compared to gpt-4o?

Image credits: gptforwork.com

# LLM Inference

Calculate by:		Input tokens:	Output tokens:	Number of API calls:	
<div>TokensWordsCharacters</div>		100	500	10000	
Provider	Model	Input price for 1M tokens	Output price for 1M tokens	Price per API call	Total price
 OpenAI	gpt-4o	\$5.00	\$15.00	\$0.0080	\$80.00
 OpenAI	gpt-4o-mini	\$0.15	\$0.60	\$0.0003	\$3.15

~~Why is gpt-4o-mini so cheap when compared to gpt-4o?~~

How can we deploy LLMs in a cost-effective manner while maintaining high performance?

Image credits: gptforwork.com

# Cost Effective Inference

1. Model Compression (lossy)

2. Efficient Engineering (lossless)



# Cost Effective Inference

1. Model Compression (lossy)

2. Efficient Engineering (lossless)

# Cost Effective Inference

## 1. Model Compression (lossy)

- 1. Quantization

- 2. Pruning

- 3. Distillation

## 2. Efficient Engineering (lossless)

# Cost Effective Inference

## 1. Model Compression (lossy)


1. **Quantization:** keep the model the same but reduce the number of bits
2. **Pruning:** remove parts of a model while retaining performance
3. **Distillation:** train a smaller model to imitate the bigger model

## 2. Efficient Engineering (lossless)

# Model Compression

1. **Quantization:** keep the model the same but reduce the number of bits
2. **Pruning:** remove parts of a model while retaining performance
3. **Distillation:** train a smaller model to imitate the bigger model

# Quantization: Problem with LLMs



- LLMs have billions of parameters which are expensive to store
- During inference, activations are created as a product of the input and the weights, which similarly are expensive to store
- The goal is to represent billions of values as efficiently as possible

Image credits: [Maarten Grootendorst](#)

# Quantization: Numerical Values Representation




Image credits: [Maarten Grootendorst](#)

# Quantization: Numerical Values Representation

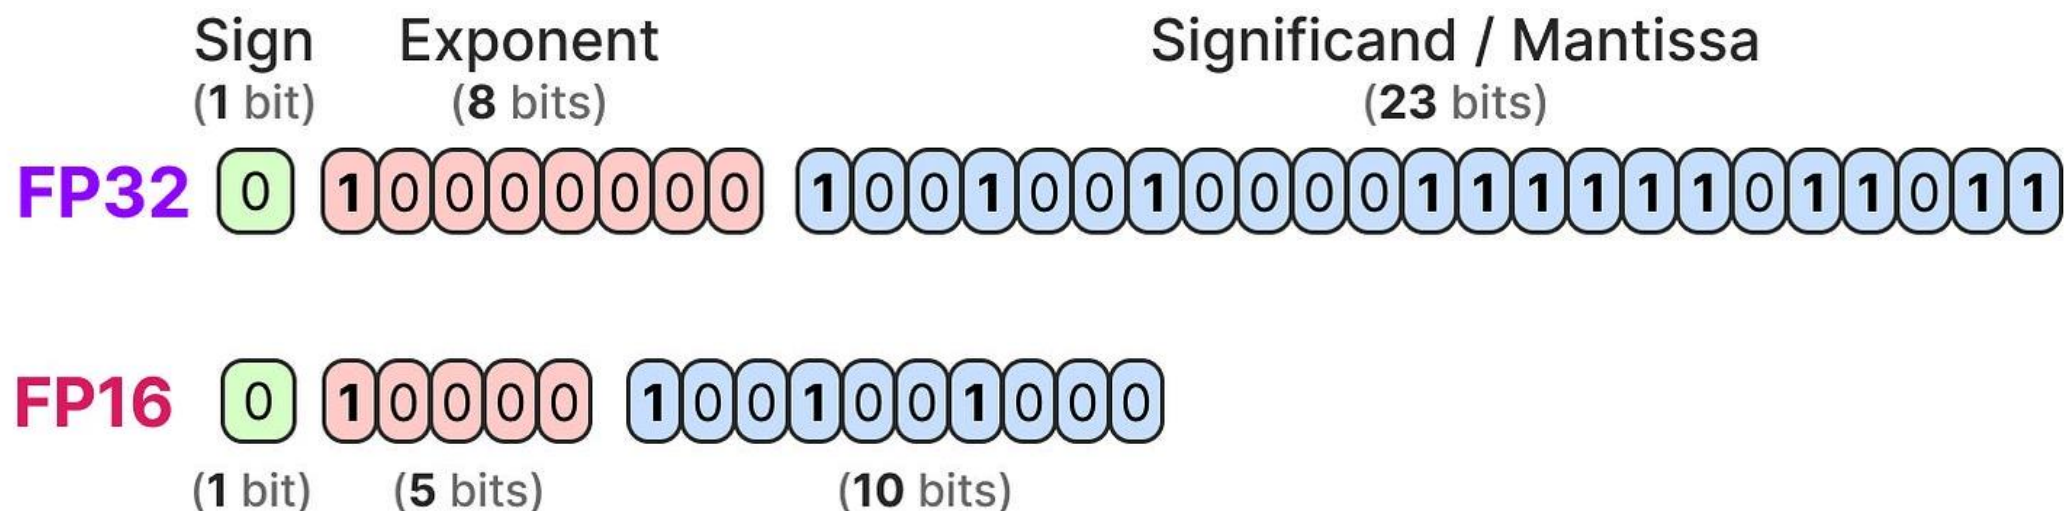


Image credits: [Maarten Grootendorst](#)

# Quantization: Numerical Values Representation




Image credits: [Maarten Grootendorst](#)



# Quantizing FP32 to INT8





Image credits: [Maarten Grootendorst](#)

# Quantizing FP32 to INT8



$$s = \frac{2^{b-1} - 1}{\alpha} \quad (\text{scale factor})$$

$$x_{\text{quantized}} = \text{round}(s \cdot x) \quad (\text{quantization})$$

$$s = \frac{127}{10.8} = 11.76 \quad (\text{scale factor})$$

$$x_{\text{quantized}} = \text{round}(11.76 \cdot \text{FP32\_value}) \quad (\text{quantization})$$

Image credits: [Maarten Grootendorst](#)

# Dequantizing INT8 to FP32



$$s = \frac{2^{b-1} - 1}{\alpha} \quad \text{(scale factor)}$$

$$x_{\text{quantized}} = \text{round}(s \cdot x) \quad \text{(quantization)}$$

$$x_{\text{dequantized}} = \frac{\text{[8 quantized values]}}{s} \quad \text{(dequantize)}$$

Image credits: [Maarten Grootendorst](#)

# Dequantizing INT8 to FP32



$$s = \frac{2^{b-1} - 1}{\alpha} \quad \text{(scale factor)}$$

$$x_{\text{quantized}} = \text{round}(s \cdot x) \quad \text{(quantization)}$$

$$x_{\text{dequantized}} = \frac{\text{[8-bit integer]}}{s} \quad \text{(dequantize)}$$

Image credits: [Maarten Grootendorst](#)

# Model Compression

1. **Quantization:** keep the model the same but reduce the number of bits

1. Post Training Quantization
2. Quantization Aware Training

2. **Pruning:** remove parts of a model while retaining performance

3. **Distillation:** train a smaller model to imitate the bigger model


# Post Training Quantization (PTQ)

- Reduce the model size without altering the LLM architecture and without **retraining**
- Weights and biases are constants. Easy to compute the scale factor(s).
- Model input and activations are variable. Use a calibration dataset to compute the scale factor(s).

# Post Training Quantization (PTQ)

## 8-bit Vector-wise Quantization

(1) Find vector-wise constants:  $C_W$  &  $C_X$



(2) Quantize

$$X_{F16} * (127/C_X) = X_{I8}$$

$$W_{F16} * (127/C_W) = W_{I8}$$

(3) Int8 Matmul

$$X_{I8} W_{I8} = Out_{I32}$$

(4) Dequantize

$$\frac{Out_{I32} * (C_X \otimes C_W)}{127 * 127} = Out_{F16}$$

Image credits: [Dettmers et al., 2022](#)

# Post Training Quantization (PTQ)

## Technical Specifications

	H100 SXM	H100 NVL
<b>FP64</b>	34 teraFLOPS	30 teraFLOPS
<b>FP64 Tensor Core</b>	67 teraFLOPS	60 teraFLOPS
<b>FP32</b>	67 teraFLOPS	60 teraFLOPS
<b>TF32 Tensor Core*</b>	989 teraFLOPS	835 teraFLOPS
<b>BFLOAT16 Tensor Core*</b>	1,979 teraFLOPS	1,671 teraFLOPS
<b>FP16 Tensor Core*</b>	1,979 teraFLOPS	1,671 teraFLOPS
<b>FP8 Tensor Core*</b>	3,958 teraFLOPS	3,341 teraFLOPS
<b>INT8 Tensor Core*</b>	3,958 TOPS	3,341 TOPS
<b>GPU Memory</b>	80GB	94GB
<b>GPU Memory Bandwidth</b>	3.35TB/s	3.9TB/s

Datasheet




## NVIDIA H100 Tensor Core GPU

Extraordinary performance, scalability,  
and security for every data center.

Image credits: nvidia.com



# PTQ: LLM.int8() [[Dettmers et al., 2022](#)]



- regular quantization retains performance at scales up to 2.7B parameters
- once systematic outliers occur at a scale of 6.7B parameters, regular quantization methods fail
- Irrespective of the scale, LLM.int8() maintains 16-bit accuracy

Image credits: [Dettmers et al., 2022](#)

# PTQ: LLM.int8()




Image credits: [Maarten Grootendorst](#)

# PTQ: LLM.int8()

## LLM.int8()




Image credits: [Dettmers et al., 2022](#)

# Model Compression

**1. Quantization:** keep the model the same but reduce the number of bits

- 1. Post Training Quantization
- 2. Quantization Aware Training

**2. Pruning:** remove parts of a model while retaining performance

**3. Distillation:** train a smaller model to imitate the bigger model

# QLoRA [[Dettmers et al. 2023](#)]

- Average memory requirements of finetuning a 65B parameter model is  $>780\text{GB}$
- QLoRA reduces the memory requirement to  $<48\text{GB}$  without degrading the predictive performance

# QLoRA [[Dettmers et al. 2023](#)]

1.4-bit NormalFloat (NF4) Quantization

2. Double Quantization

3. Paged Optimizers

4. LoRA

# QLoRA [[Dettmers et al. 2023](#)]

1. NF4 Quantization
2. Double Quantization
3. Paged Optimizers
4. LoRA





Image credits: [Shaw Talebi](#)

# QLoRA [[Dettmers et al. 2023](#)]

1. NF4 Quantization
2. Double Quantization
3. Paged Optimizers
4. LoRA

Double Quantization is the process of quantizing the quantization constants for additional memory savings

Stored as FP32


$$S = \frac{2^{b-1} - 1}{\alpha}$$



# QLoRA [[Dettmers et al. 2023](#)]

1. NF4 Quantization
2. Double Quantization
3. Paged Optimizers
4. LoRA




Image credits: [Shaw Talebi](#)

# QLoRA [[Dettmers et al. 2023](#)]

1. NF4 Quantization
2. Double Quantization
3. Paged Optimizers
4. LoRA




Image credits: [[Hu et al., 2022](#)]

# QLoRA [[Dettmers et al. 2023](#)]


Image credits: [Dettmers et al. 2023](#), exxactcorp.com

$$\mathbf{Y} = \mathbf{XW} + s\mathbf{XL}_1\mathbf{L}_2$$

# QLoRA [[Dettmers et al. 2023](#)]

Image credits: [Dettmers et al. 2023](#), exxactcorp.com

$$\mathbf{Y} = \mathbf{XW} + s\mathbf{XL}_1\mathbf{L}_2$$




# QLoRA [Dettmers et al. 2023]

Image credits: [Dettmers et al. 2023](#), [exxactcorp.com](#)

$$\mathbf{Y} = \mathbf{XW} + s\mathbf{XL}_1\mathbf{L}_2$$


$$\mathbf{Y}^{\text{BF16}} = \mathbf{X}^{\text{BF16}} \text{doubleDequant}(c_1^{\text{FP32}}, c_2^{\text{k-bit}}, \mathbf{W}^{\text{NF4}}) + \mathbf{X}^{\text{BF16}} \mathbf{L}_1^{\text{BF16}} \mathbf{L}_2^{\text{BF16}},$$

$$\text{doubleDequant}(c_1^{\text{FP32}}, c_2^{\text{k-bit}}, \mathbf{W}^{\text{k-bit}}) = \text{dequant}(\text{dequant}(c_1^{\text{FP32}}, c_2^{\text{k-bit}}), \mathbf{W}^{\text{4bit}}) = \mathbf{W}^{\text{BF16}}$$



# QLoRA [[Dettmers et al. 2023](#)]

Image credits: [Dettmers et al. 2023](#)



- NFloat data type improves the bit-for-bit accuracy gains compared to regular 4-bit Floats
- Double Quantization (DQ) only leads to minor gains, it allows for a more fine-grained control over the memory footprint

Mean zero-shot accuracy over Winogrande, HellaSwag, PiQA, Arc-Easy, and Arc-Challenge using LLaMA models with different 4-bit data types.

# QLoRA [[Dettmers et al. 2023](#)]

Image credits: [Dettmers et al. 2023](#)

LLaMA Size Dataset	Mean 5-shot MMLU Accuracy								Mean
	7B		13B		33B		65B		
	Alpaca	FLAN v2	Alpaca	FLAN v2	Alpaca	FLAN v2	Alpaca	FLAN v2	
BFloat16	38.4	45.6	47.2	50.6	57.7	60.5	61.8	62.5	53.0
Float4	37.2	44.0	47.3	50.0	55.9	58.5	61.3	63.3	52.2
NFloat4 + DQ	39.0	44.5	47.5	50.7	57.3	59.2	61.8	63.9	53.1

Mean 5-shot MMLU test accuracy for LLaMA models finetuned with adapters on Alpaca and FLAN v2 for different data types.


# Model Compression

1. **Quantization:** keep the model the same but reduce the number of bits
2. **Pruning:** remove parts of a model while retaining performance
3. **Distillation:** train a smaller model to imitate the bigger model



# Pruning

## Unstructured Pruning



## Structured Pruning





Image credits: neuralmagic.com

# Magnitude Pruning [[Han et al. 2015](#), [See et al. 2016](#)]



- prune weights with smallest absolute value
- prunes 40% of the weights with negligible performance loss
- by adding a retraining phase after pruning, we can prune 80% with no performance loss

Image credits: [See et al. 2016](#)

# Wanda [Sun et al. 2023]

## Magnitude Pruning

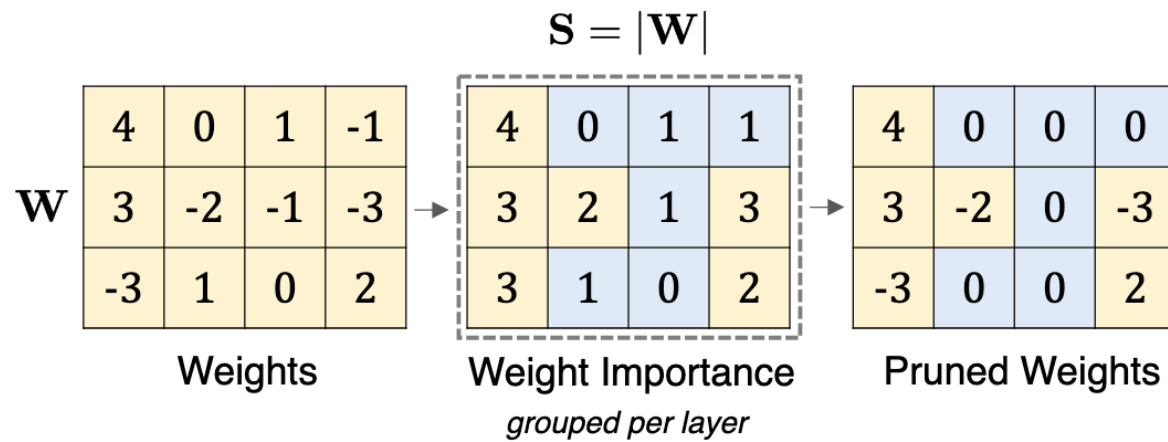
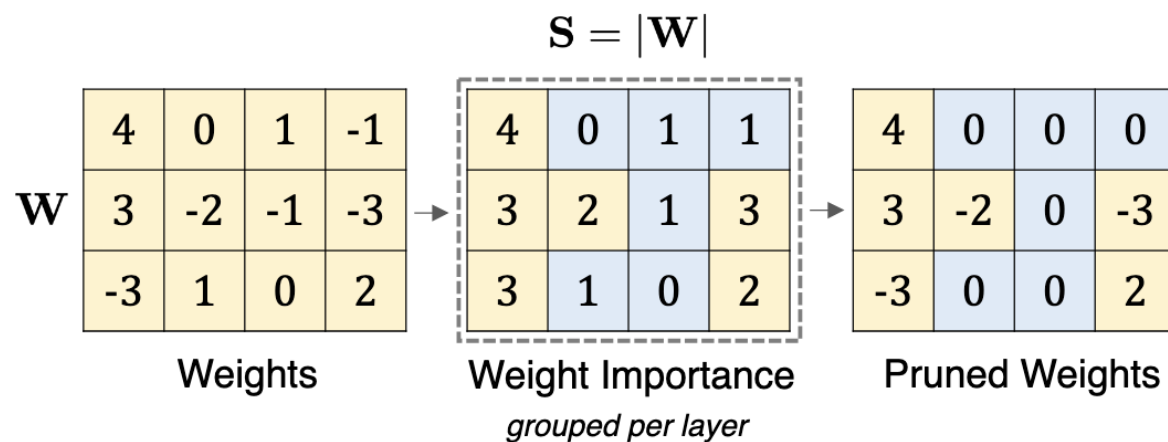


Image credits: [Sun et al. 2023](#)

# Wanda [Sun et al. 2023]

## Magnitude Pruning



## Wanda

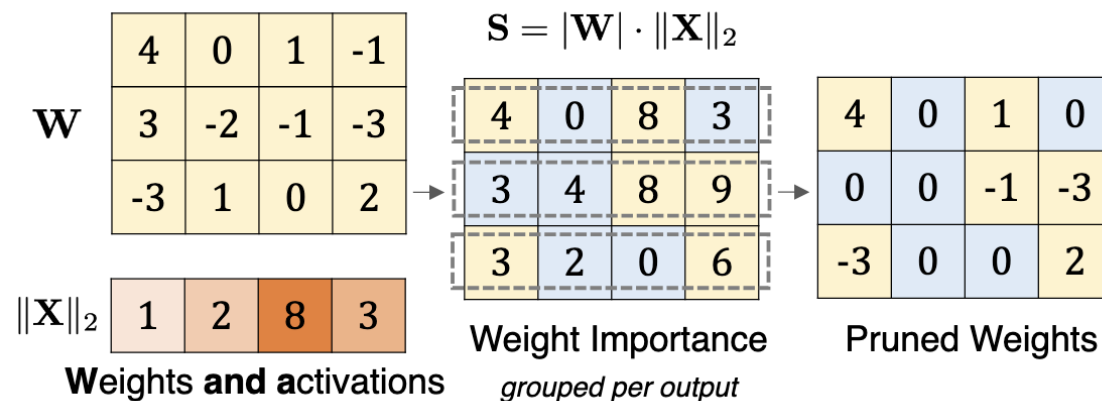


Image credits: [Sun et al. 2023](#)

# Unstructured Pruning

Unstructured pruning can work only if the hardware supports.





Image credits: nvidia.com

# Structured Pruning




- NVIDIA A100 GPU supports fine-grained structured sparsity to its Tensor Cores
- Sparse Tensor Cores accelerate a 2:4 sparsity pattern.

Image credits: nvidia.com

# Structured Pruning

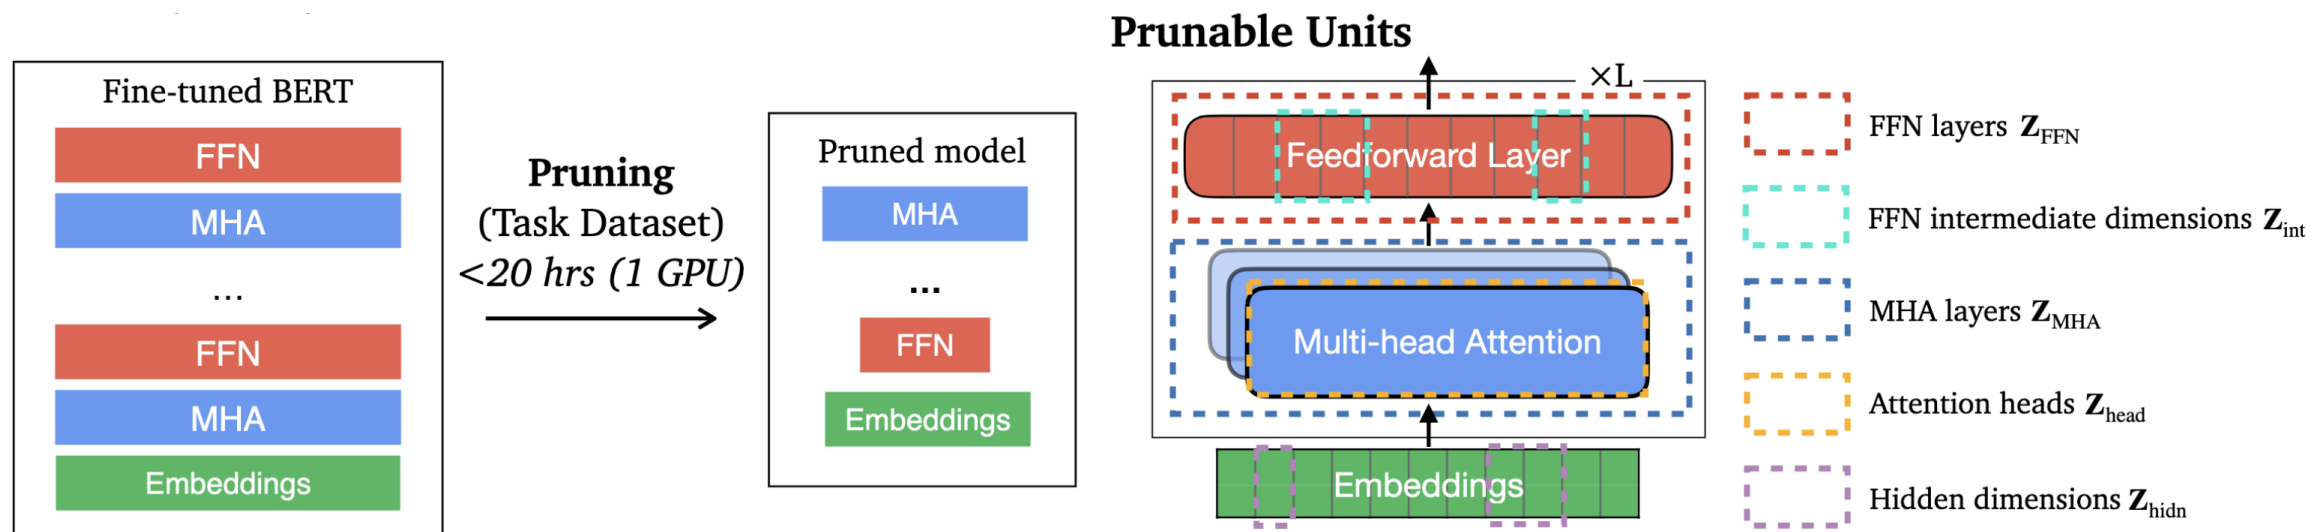
Input Operands	Accumulator	Dense TOPS	vs. FFMA	Sparse TOPS	vs. FFMA
FP32	FP32	19.5	–	–	–
TF32	FP32	156	8X	<b>312</b>	<b>16X</b>
FP16	FP32	312	16X	<b>624</b>	<b>32X</b>
BF16	FP32	312	16X	<b>624</b>	<b>32X</b>



- NVIDIA A100 GPU supports fine-grained structured sparsity to its Tensor Cores
- Sparse Tensor Cores accelerate a 2:4 sparsity pattern.

Image credits: nvidia.com

# Structured Pruning [Xia et al. 2022]

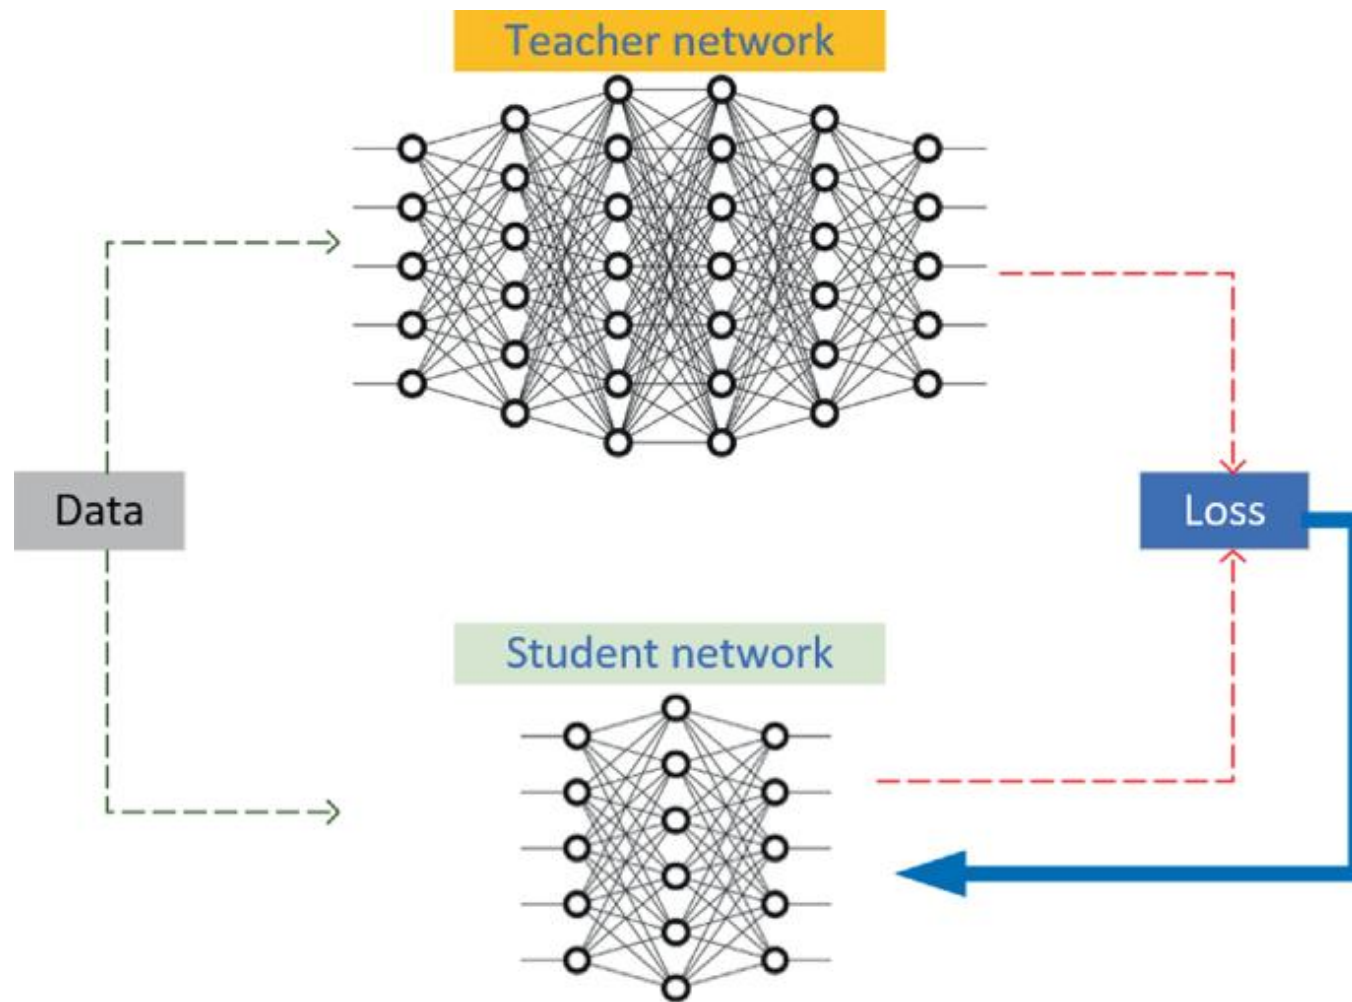




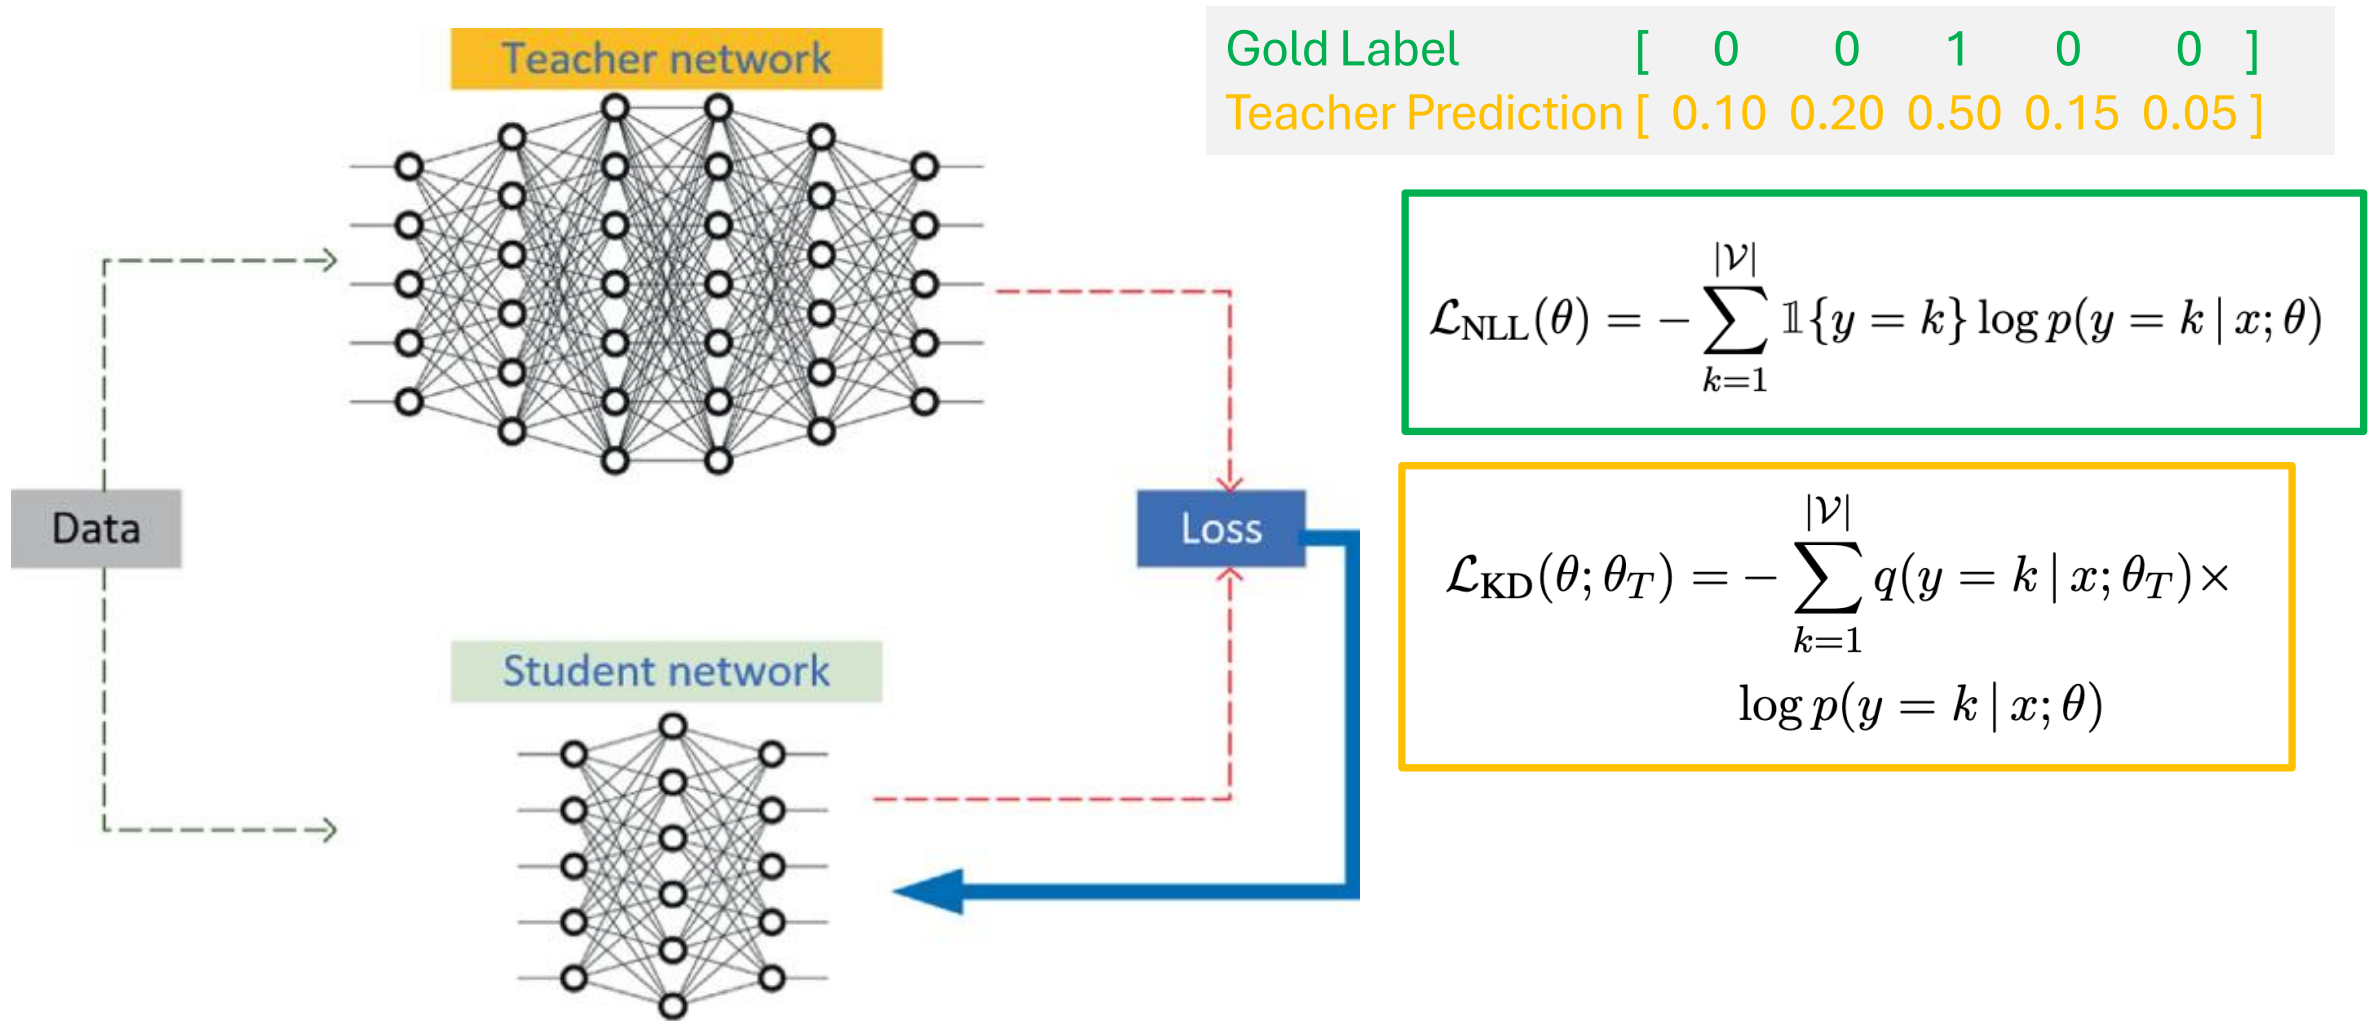
# Model Compression

1. **Quantization:** keep the model the same but reduce the number of bits
2. **Pruning:** remove parts of a model while retaining performance
3. **Distillation:** train a smaller model to imitate the bigger model

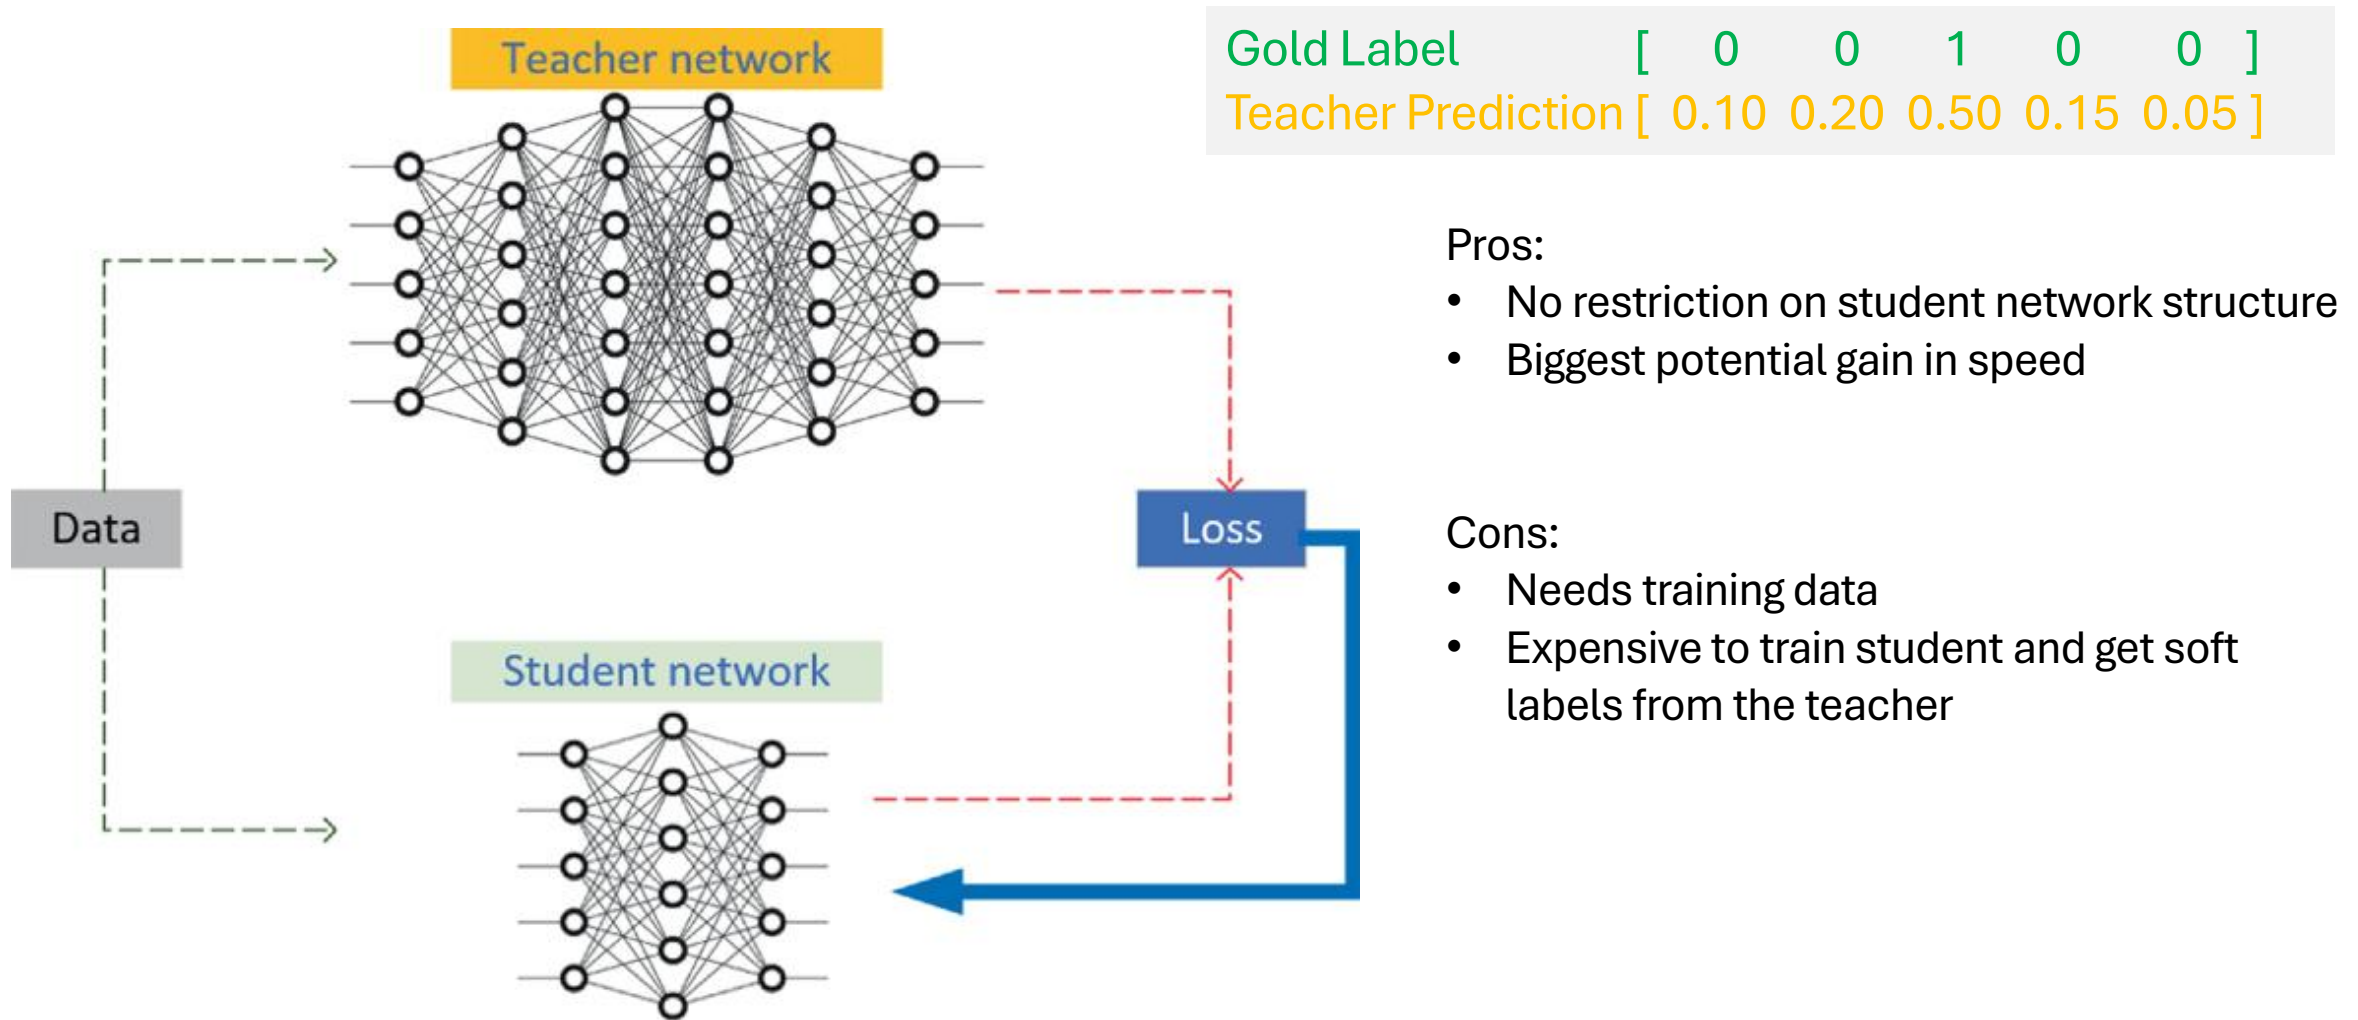
# Distillation [Hinton et al 2015]



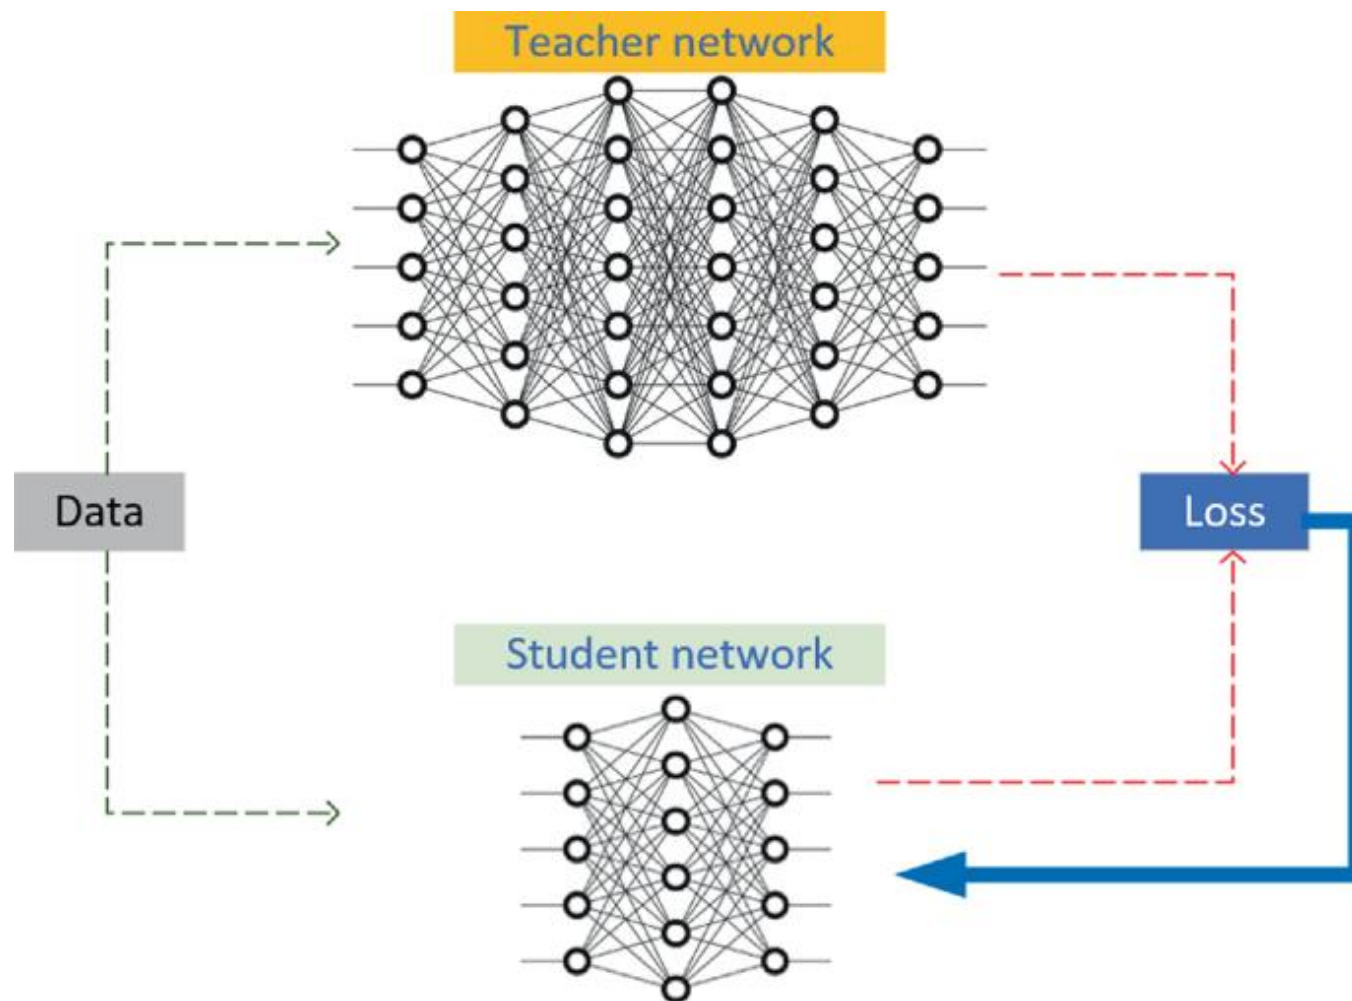
# Distillation [Hinton et al 2015]



# Distillation [Hinton et al 2015]



# Distillation [Hinton et al 2015]



Gold Label	[	0	0	1	0	0	]
Soft Target	[	0.90	0.01	0.05	0.01	0.03	]
Hard Target	[	1.	0.	0	0	0	]

# Sequence Level Distillation [Kim et al. 2016]

$$\mathcal{L}_{\text{KD}}(\theta; \theta_T) = - \sum_{k=1}^{|\mathcal{V}|} q(y = k | x; \theta_T) \times \log p(y = k | x; \theta)$$

# Sequence Level Distillation [Kim et al. 2016]

## 1. Word-Level Knowledge Distillation

$$\mathcal{L}_{\text{WORD-KD}} = - \sum_{j=1}^J \sum_{k=1}^{|\mathcal{V}|} q(t_j = k | \mathbf{s}, \mathbf{t}_{<j}) \times \log p(t_j = k | \mathbf{s}, \mathbf{t}_{<j})$$

$$\mathcal{L}_{\text{KD}}(\theta; \theta_T) = - \sum_{k=1}^{|\mathcal{V}|} q(y = k | x; \theta_T) \times \log p(y = k | x; \theta)$$

# Sequence Level Distillation [Kim et al. 2016]

## 1. Word-Level Knowledge Distillation

$$\mathcal{L}_{\text{WORD-KD}} = - \sum_{j=1}^J \sum_{k=1}^{|\mathcal{V}|} q(t_j = k | \mathbf{s}, \mathbf{t}_{<j}) \times \log p(t_j = k | \mathbf{s}, \mathbf{t}_{<j})$$

$$\mathcal{L}_{\text{KD}}(\theta; \theta_T) = - \sum_{k=1}^{|\mathcal{V}|} q(y = k | x; \theta_T) \times \log p(y = k | x; \theta)$$

## 2. Sequence-Level Knowledge Distillation

$$\begin{aligned} \mathcal{L}_{\text{SEQ-KD}} &= - \sum_{\mathbf{t} \in \mathcal{T}} q(\mathbf{t} | \mathbf{s}) \log p(\mathbf{t} | \mathbf{s}) \\ &\approx - \sum_{\mathbf{t} \in \mathcal{T}} \mathbb{1}\{\mathbf{t} = \hat{\mathbf{y}}\} \log p(\mathbf{t} | \mathbf{s}) \\ &= - \log p(\mathbf{t} = \hat{\mathbf{y}} | \mathbf{s}) \end{aligned}$$



# Sequence Level Distillation [Kim et al. 2016]

## 1. Word-Level Knowledge Distillation

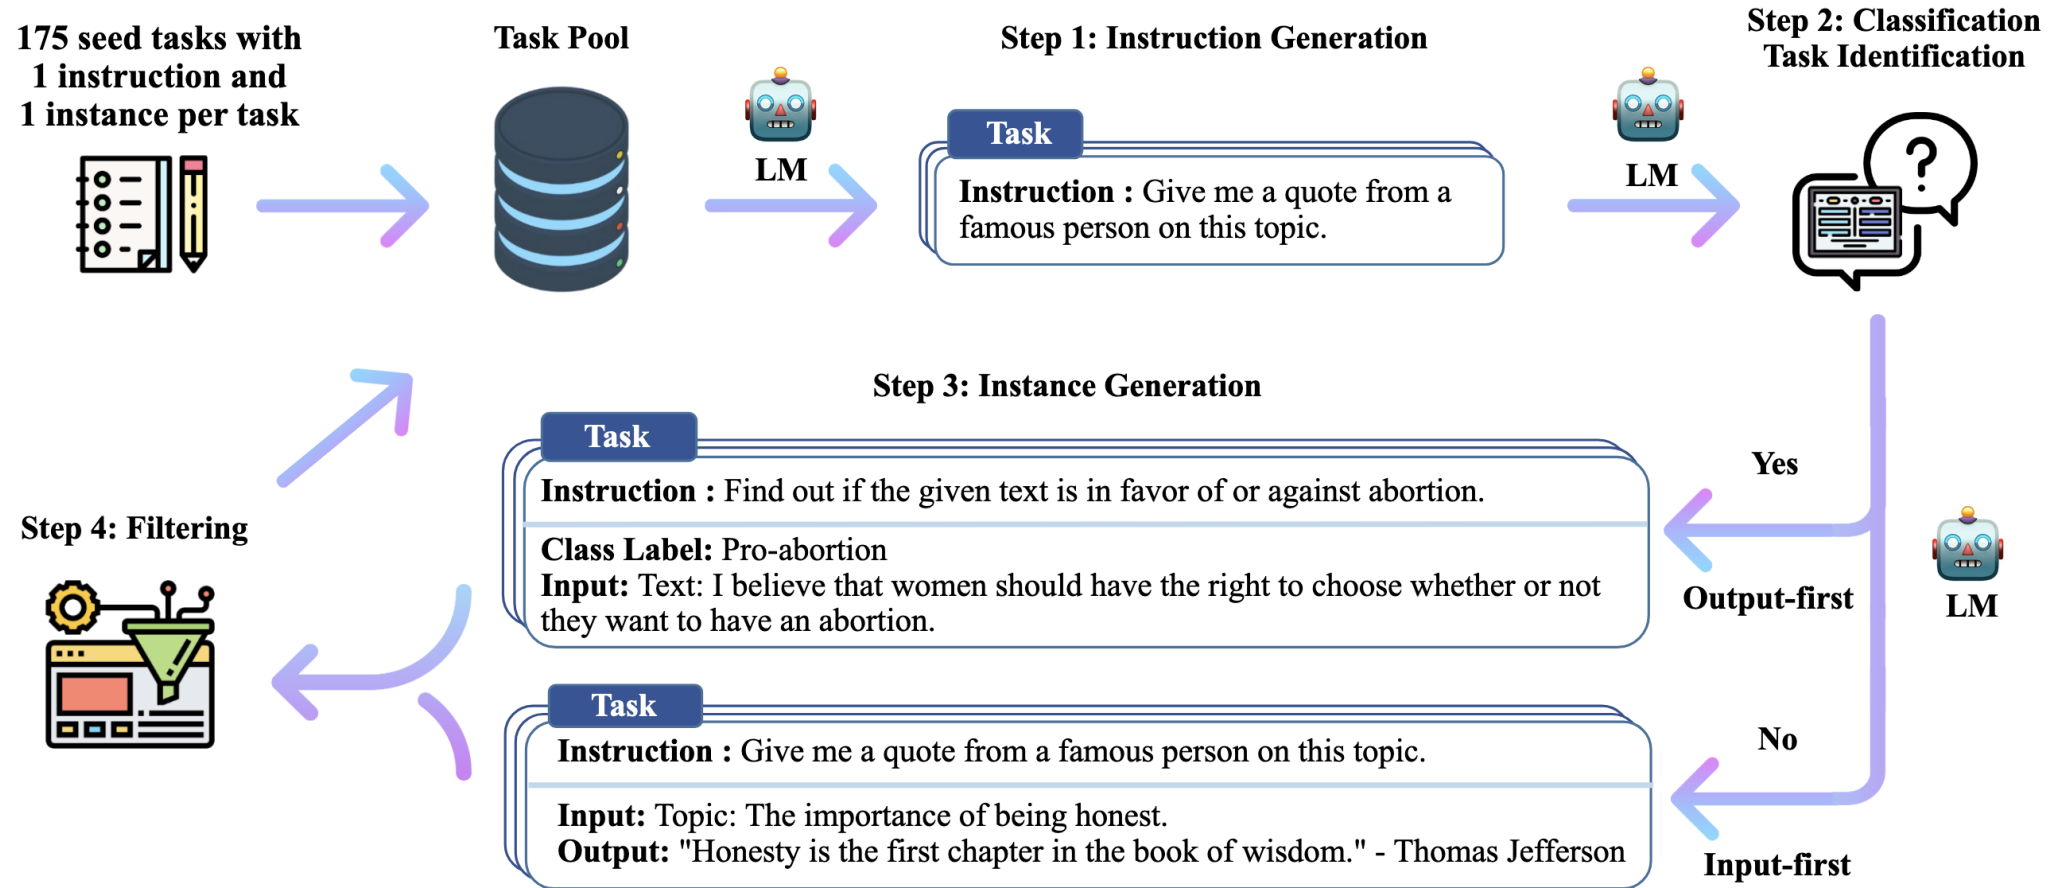
$$\mathcal{L}_{\text{WORD-KD}} = - \sum_{j=1}^J \sum_{k=1}^{|\mathcal{V}|} q(t_j = k | \mathbf{s}, \mathbf{t}_{<j}) \times \log p(t_j = k | \mathbf{s}, \mathbf{t}_{<j})$$

$$\mathcal{L}_{\text{KD}}(\theta; \theta_T) = - \sum_{k=1}^{|\mathcal{V}|} q(y = k | x; \theta_T) \times \log p(y = k | x; \theta)$$

## 2. Sequence-Level Knowledge Distillation

$$\mathcal{L}_{\text{SEQ-KD}} = - \sum_{\mathbf{t} \in \mathcal{T}} q(\mathbf{t} | \mathbf{s}) \log p(\mathbf{t} | \mathbf{s})$$

# Self-Instruct [Wang et al. 2023]



# Model Compression

1. **Quantization:** keep the model the same but reduce the number of bits
2. **Pruning:** remove parts of a model while retaining performance
3. **Distillation:** train a smaller model to imitate the bigger model